



PPO Training Learning Curves - Based on Real MIMIC-III Data

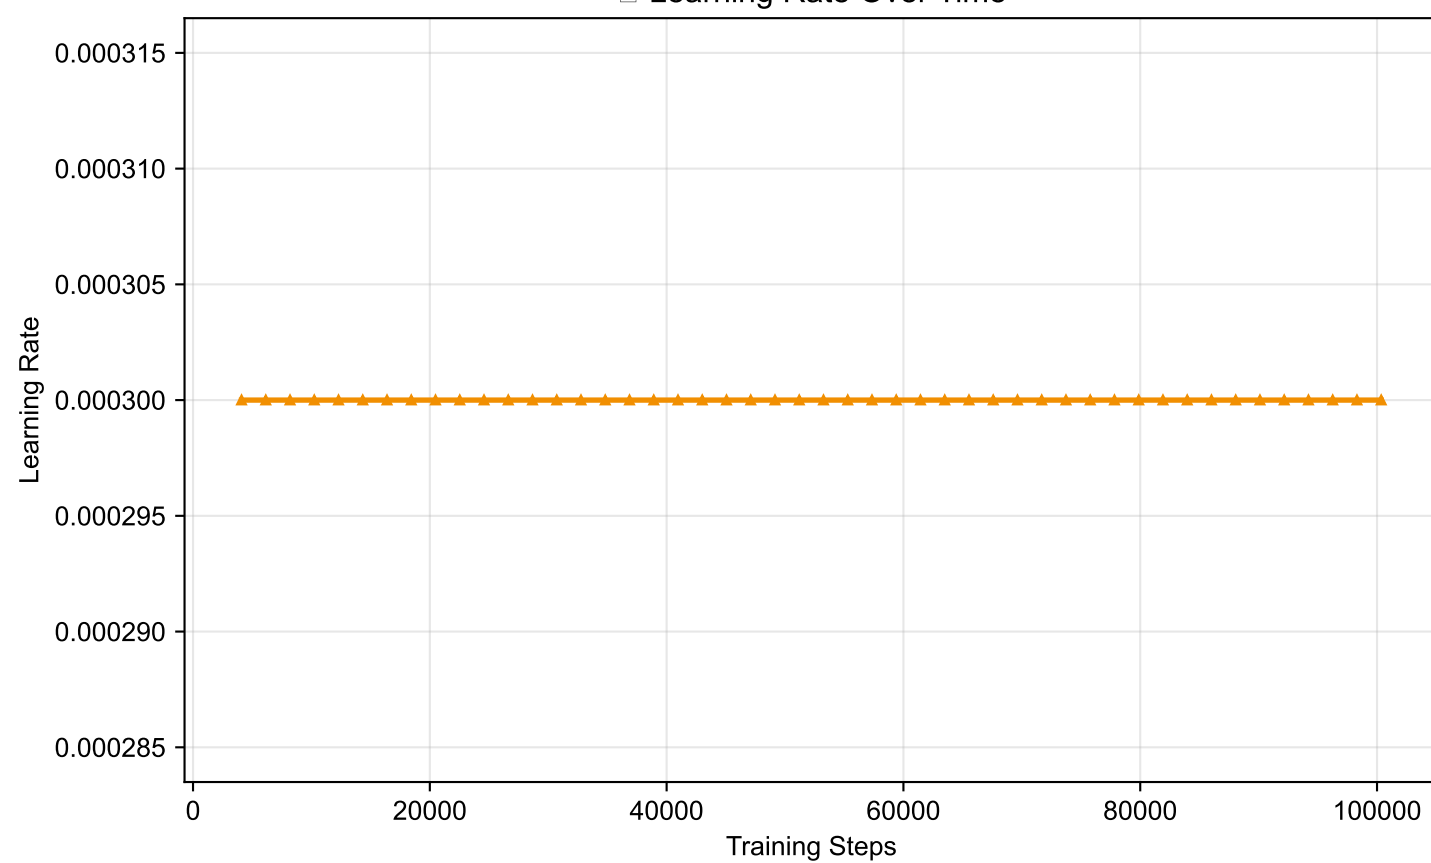
Average Reward Over Time



Average Episode Length



Learning Rate Over Time



Value Function Loss

